



Pokemon Data Ingestion

This document explains the steps to ingest the Pokemon data that will be used by the Frontend and Backend workshops in the AWE-AUB context.

In this tutorial, you will learn how to:

[Pokemon data ingestion](#)

[Step 1 - Select London \(eu-west-2\) as your region](#)

[Step 2 - Create a DynamoDB table](#)

[Step 3 - Ingest one Pokemon manually to a DynamoDB table](#)

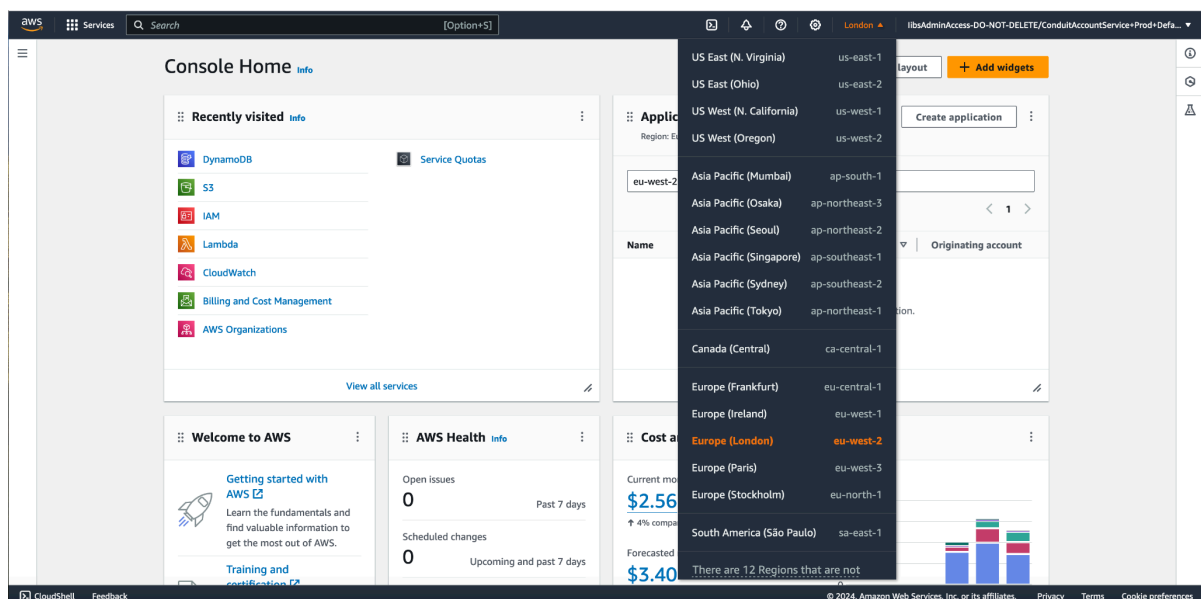
[Step 4 - Create a DynamoDB data ingestion lambda](#)

[Step 5 - Invoke the data ingestion lambda](#)

[Step 6 - Verify that your Pokemons DynamoDB was correctly populated](#)

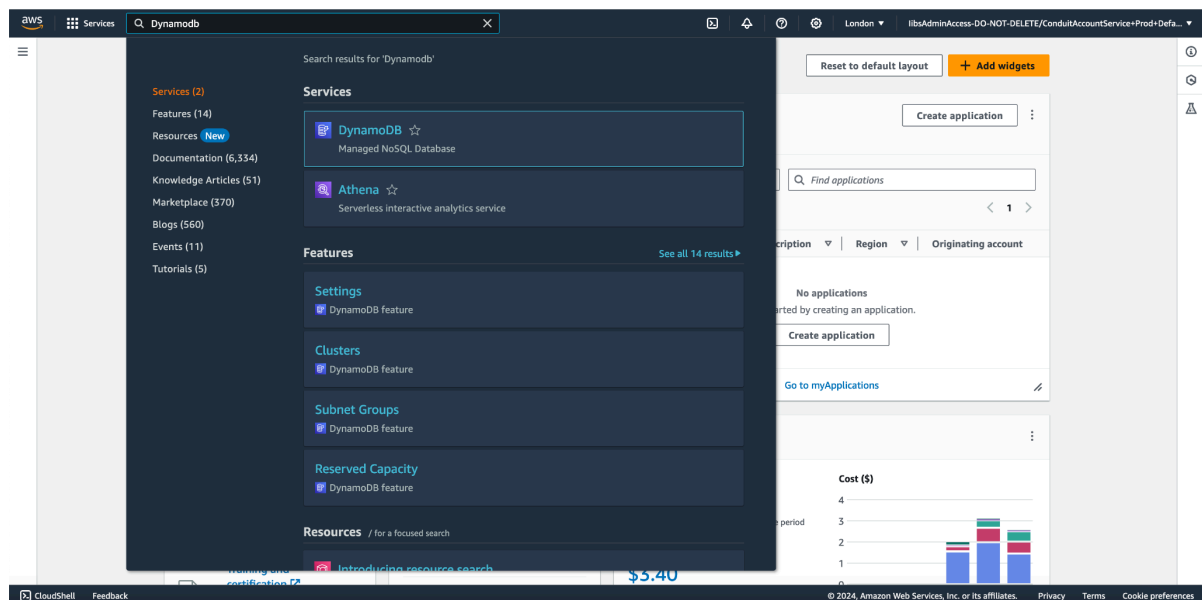
Step 1 - Select London (eu-west-2) as your region

This tutorial expects that the resources created are in the London (eu-west-2) AWS region. Before starting any actual development, select this region on the AWS console.

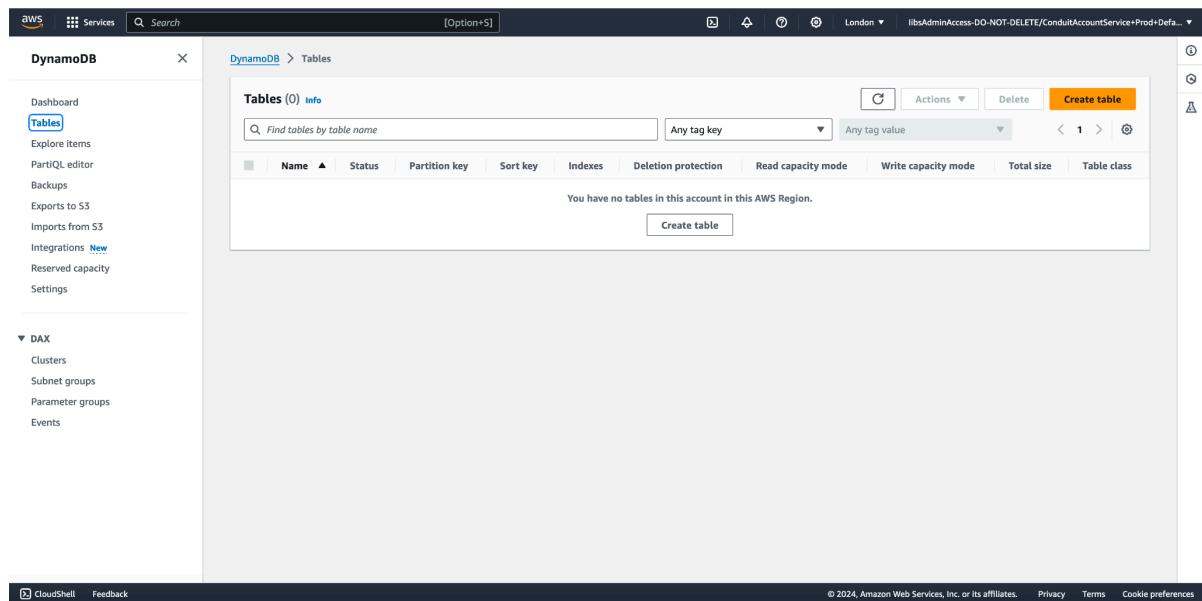


Step 2 - Create a DynamoDB table

Amazon DynamoDB is a **serverless, NoSQL database service that enables you to develop modern applications at any scale**. As a serverless database, you only pay for what you use and DynamoDB has no extra latency for read requests when these requests are made very sporadically, no version upgrades, no maintenance windows, no need to perform software updates, and no downtime maintenance. The reason why we chose DynamoDB is because the records you add to this table can have different formats, which will be very helpful in the early phases of your project as the data model can change very quickly. Imagine that to launch a new feature to your project you need to start storing a new field in your records, with DynamoDB it's very convenient to manipulate new records or to patch the old ones with new fields without having to perform migrations from one table to another. To start, on the AWS console, search for **DynamoDB**:



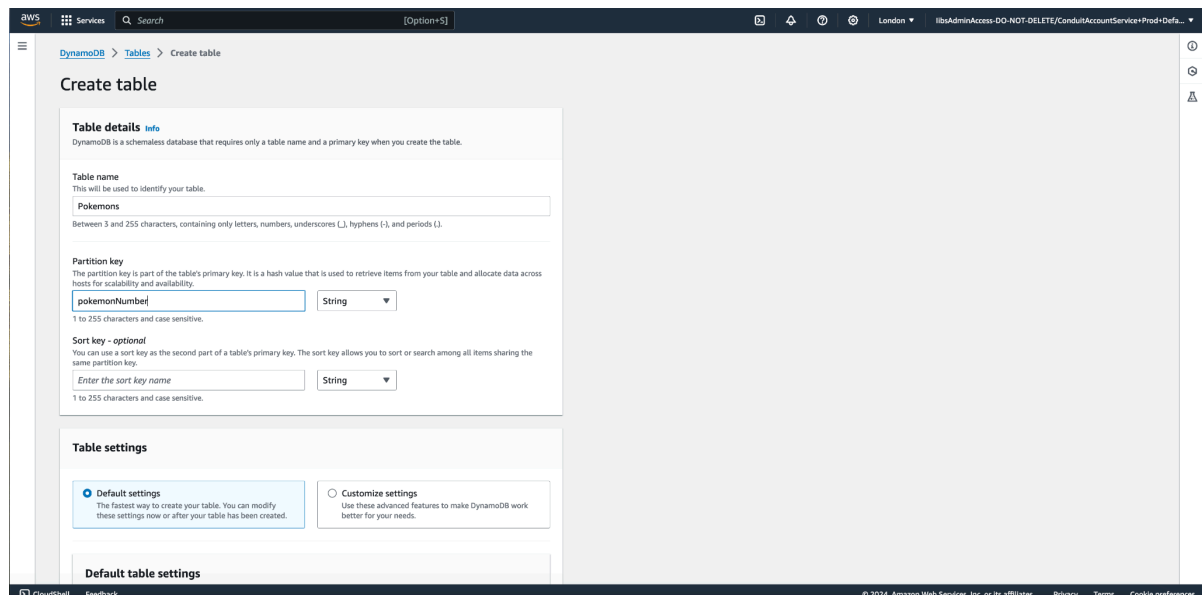
On the left side, select **Tables**. Then **Create table** on the right.



Make sure to create a table with the following inputs, **future steps of this series of workshops have a hard dependency on these names:**

- Table name: **Pokemons**
- Partition key: **pokemonNumber** and the type should be: **Number**
- Everything else can be as is by default

Click create table at the bottom of the page and wait for a couple of seconds for your table to be created:



The table creation will take a while and you should see a spinning widget.



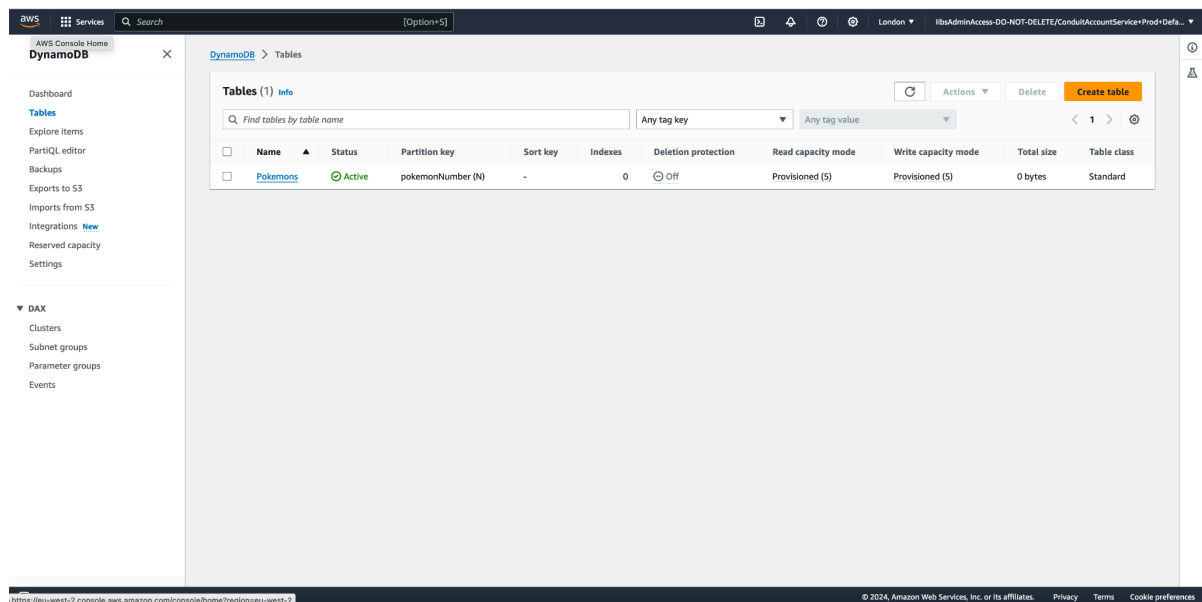
After a couple of seconds, you should see a successfully created table widget.

🟢 The Pokemons table was created successfully.

That table will contain all Pokemons that will later be used through the rest of this Pokemon workshop.

Step 3 - Ingest one Pokemon manually to a DynamoDB table

Open your table by clicking its name (you might need to refresh your screen to reload the dashboard with the created tables).



On the actions dropdown, select **Create item**:

Amazon Web Services console screenshot showing the DynamoDB console for the 'Pokemons' table. The left sidebar shows navigation options like Dashboard, Tables, Explore items, PartiQL editor, Backups, Exports to S3, Imports from S3, Integrations, Reserved capacity, and Settings. The main content area displays the 'Pokemons' table overview, including a warning about protecting the table from accidental writes and deletes, general information (Partition key: pokemonNumber, Sort key: -, Capacity mode: Provisioned), and items summary (Item count: 0, Table size: 0 bytes, Average item size: 0 bytes). The 'Actions' menu is open, showing options like Edit capacity, Update table class, Delete table, Create item, Create index, Create replica, Export to S3, Turn on TTL, Manage tags, and Create access control policy. The 'JSON view' option is highlighted.

Select the **JSON view** on the top right and paste the following content to manually ingest a Bulbassaur to your database:

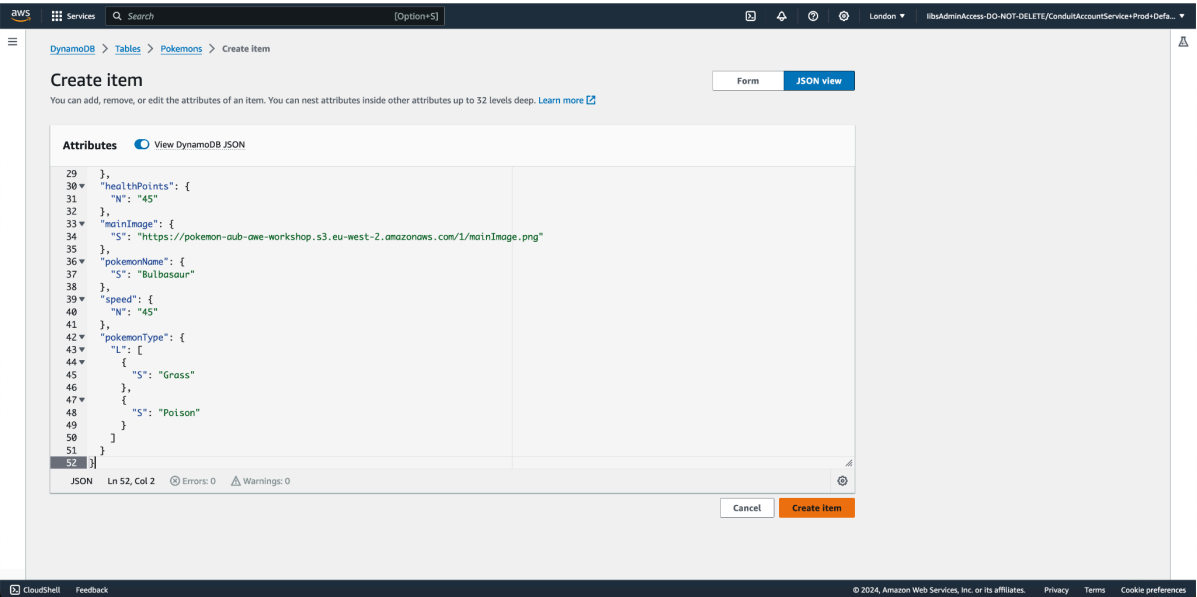
Amazon Web Services console screenshot showing the 'Create item' page for the 'Pokemons' table. The page has a 'Form' and 'JSON view' tab, with 'JSON view' selected. The 'Attributes' section shows a JSON object:

```
{  "pokemonNumber": {    "N": "0"  }}
```

. The bottom of the page has 'Cancel' and 'Create item' buttons.

```
{
  "pokemonNumber": {
    "N": "1"
  },
  "attack": {
    "N": "49"
  },
  "defense": {
    "N": "49"
  },
  "devolution": {
    "S": ""
  },
  "evolution": {
    "S": "Ivysaur"
  },
  "evolutionFamily": {
    "L": [
      {
        "S": "Bulbasaur"
      },
      {
        "S": "Ivysaur"
      },
      {
        "S": "Venusaur"
      }
    ]
  },
  "healthPoints": {
    "N": "45"
  },
  "mainImage": {
    "S":
"https://pokemon-aub-awe-workshop.s3.eu-west-2.amazonaws.com/1/mainImage.png"
  },
  "pokemonName": {
    "S": "Bulbasaur"
  },
  "speed": {
    "N": "45"
  },
  "pokemonType": {
    "L": [
      {
```

```
    "S": "Grass"
  },
  {
    "S": "Poison"
  }
]
}
```



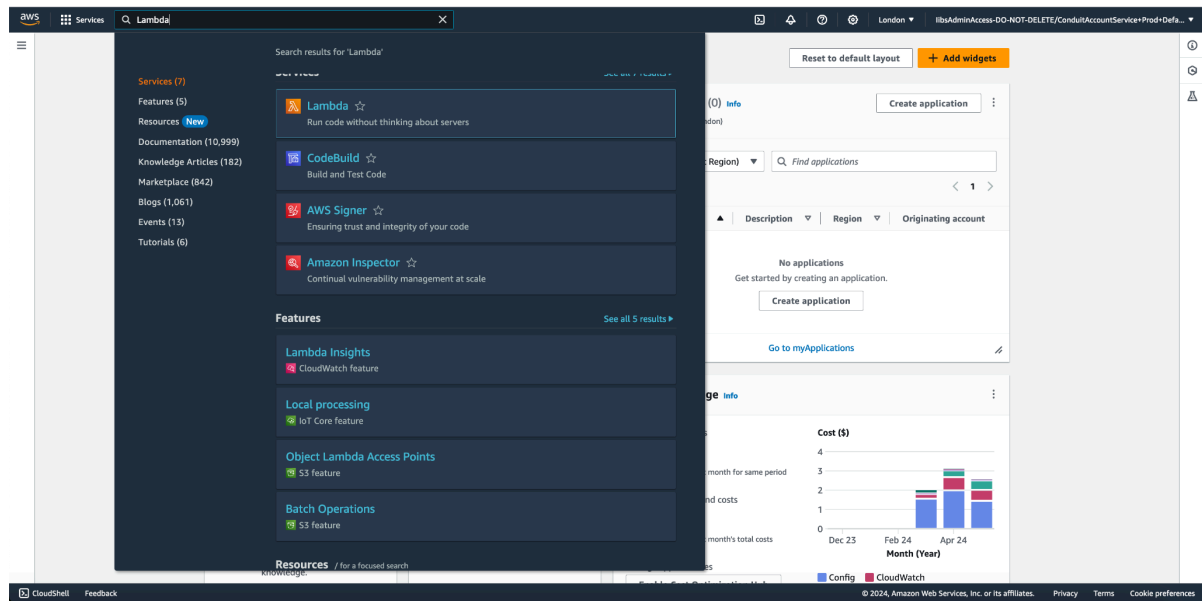
After pasting the content, click the [Create item](#) button. To verify that one item is added to the table click on [Explore table items](#) on the top right. Congratulations, you have added a Bulbasaur to your table.



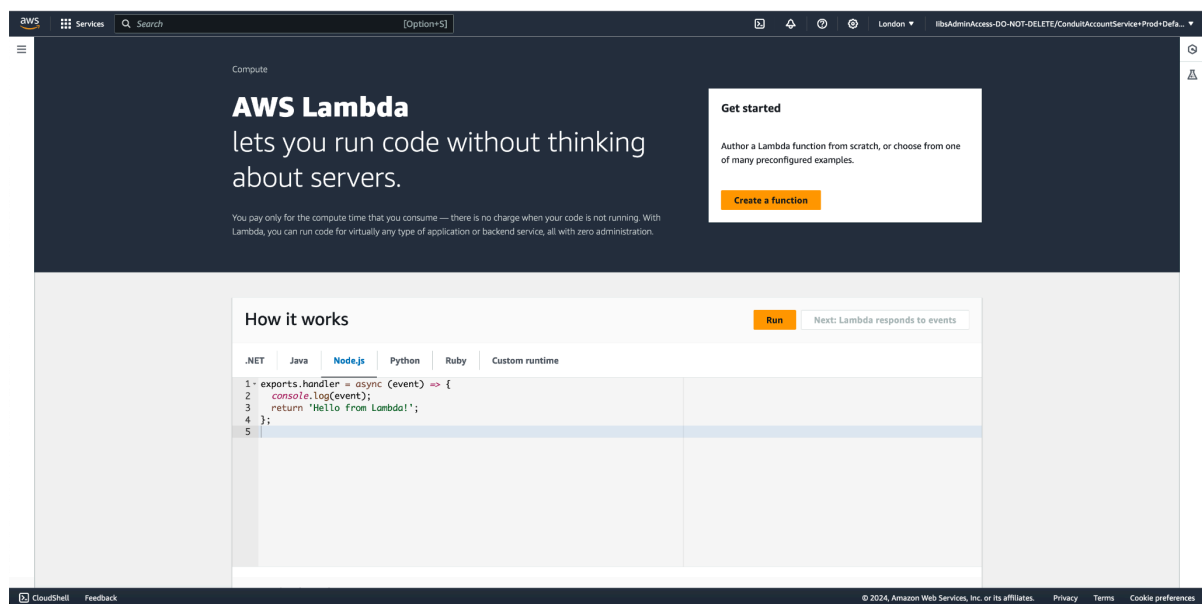
Step 4 - Create a DynamoDB data ingestion lambda

AWS Lambda is a **serverless, event-driven compute service that lets you run code for virtually any type of application or backend service without provisioning or managing servers**. Now that we have created a table and tested the manual ingestion process to get familiar with the AWS ecosystem and the DynamoDB table we created, let's create an AWS Lambda that will ingest the remaining 150 Pokemons to the DynamoDB table for us. Notice that we could have inserted manually the remaining 150 Pokemons, but that would be a tedious task, but also notice that if we had to insert 5 records to our table, creating a lambda to perform this task would take more time than just inserting the data manually ourselves. Always balance the time to automate a task vs doing it manually, and account for the number of times you will need to do it manually again in the future.

To get started with your lambda creation, on the AWS console, search for [Lambda](#):



Select **Create function** on the top right:



Select **Author from scratch** and make sure that the Runtime is **Python 3.12** and the Architecture is **x86_64**. For the name, use something like **PokemonDataIngestionLambda**.

Basic information

Function name
Enter a name that describes the purpose of your function.
PokemonDataIngestionLambda

Runtime [info](#)
Choose the language to use to write your function. Note that the console code editor supports only Node.js, Python, and Ruby.
Python 3.12

Architecture [info](#)
Choose the instruction set architecture you want for your function code.
☒ x86_64
☐ arm64

Permissions [info](#)
By default, Lambda will create an execution role with permissions to upload logs to Amazon CloudWatch Logs. You can customize this default role later when adding triggers.

Change default execution role

Execution role
Choose a role that defines the permissions of your function. To create a custom role, go to the [IAM console](#).

☐ Create a new role with basic Lambda permissions
☐ Use an existing role
☒ Create a new role from AWS policy templates

Role creation might take a few minutes. Please do not delete the role or edit the trust or permissions policies in this role.

Role name
Enter a name for your new role.
pokemon-data-ingestion-lambda

Policy templates - optional [info](#)
Choose one or more policy templates.
Simple microservice permissions X
DynamoDB

Expand the [Change default execution role](#) dropdown and select:

- Create a new role from AWS policy templates
- For role name, use [pokemon-data-ingestion-lambda](#)
- And select [Simple microservice permissions](#) and [Amazon s3 object read-only permissions](#) as policy template, notice that these are DynamoDB permissions.

Then proceed with [Create function](#). After your function is created, select the [Configuration](#) tab.

PokemonDataIngestionLambda

Function overview [info](#)

Diagram **Template**

General configuration [info](#) [Edit](#)

Description	Memory	Ephemeral storage
-	128 MB	512 MB
Timeout	Snapshot	
0 min 3 sec	None	

By default, the timeout of lambdas is configured to 3 seconds. However, to ingest the pokemons to the database, it takes approximately 10 seconds. Click [Edit](#) and change the [timeout](#) to 1 minute and 0 seconds and then click [Save](#).

Step 5 - Invoke the data ingestion lambda

From reading the python code below, you will notice that it's reading from an S3 bucket a csv file.

Amazon Simple Storage Service (Amazon S3) is an **object storage service** that offers industry-leading scalability, data availability, security, and performance. Customers of all sizes and industries can use Amazon S3 to store and protect any amount of data for a range of use cases, such as data lakes, websites, mobile applications, backup and restore, archive, enterprise applications, IoT devices, and big data analytics. Amazon S3 provides management features so that you can optimize, organize, and configure access to your data to meet your specific business, organizational, and compliance requirements.

Below are the two first lines of the .csv file we are reading. It's a simple csv file containing the Pokemon information we need to build our pokédex. Notice that we are converting the rows from this csv file to records in the DynamoDB table, we could have designed our application in a way where we wouldn't need the DynamoDB table and we could make reads to this csv file. Going with the DynamoDB allows us to have more flexibility on the data model and allows us for faster reads. Imagine that you want to read information from a **pokemonNumber** of 3, using DynamoDB that's an $O(1)$ operation, as the **pokemonNumber** is a key of that table. If we were using the csv file directly, that would be an $O(n)$ operation as to get the pokemon with **pokemonNumber** of 3 we would potentially need to read the whole csv file. In that case, the csv file is sorted by **pokemonNumber** so you could try to guess where it is, but that assumption of the csv file being sorted could break at any time.

```
pokemonNumber,pokemonName,pokemonType,healthPoints,attack,defense,speed,evolution,devolution,evolutionFamily,mainImage
1,Bulbasaur,"['Grass', 'Poison']",45,49,49,45,Ivysaur,,['Bulbasaur', 'Ivysaur', 'Venusaur']",https://pokemon-aub-awe-workshop.s3.eu-west-2.amazonaws.com/1/mainImage.png
2,Ivysaur,"['Grass', 'Poison']",60,62,63,60,Venusaur,Bulbasaur,['Bulbasaur', 'Ivysaur', 'Venusaur']",https://pokemon-aub-awe-workshop.s3.eu-west-2.amazonaws.com/2/mainImage.png
```

On the lambda you have created, move to the **Code** tab, and paste the following code in the **lambda_function** tab of the code editor:

```
import boto3
import ast
import io
import csv

region = 'eu-west-2'
# Initialize dynamodb resource
dynamodb = boto3.resource('dynamodb', region_name=region)
# Initialize the S3 client
s3 = boto3.client('s3')
```

```

# Specify the S3 bucket and object key of the CSV file
bucket_name = 'pokemon-aub-awe-workshop'
file_key = 'final_pokemon.csv'

# DynamoDB table name
dynamodb_table = 'Pokemons'

def read_csv_from_s3_bucket(bucket_name: str, file_key: str):
    """
    Read a CSV file from S3 bucket

    Arguments:
    - bucket_name: name of the S3 bucket to read files from
    - file_key: file path where to read from
    """
    response = s3.get_object(Bucket=bucket_name, Key=file_key)
    csv_content = response['Body'].read().decode('utf-8')

    return csv_content

def insert_dynamoDB_items(tablename: str, items: list):
    """
    Insert items to a dynamodb table.

    Arguments:
    - tablename: name of the dynamodb table where items will be inserted
    - items: list of elements that will be added to the dynamodb table
    """
    ddb = dynamodb.Table(tablename)

    for record in items:
        ddb.put_item(Item=record)

def apply_literal_eval(items: list):
    """
    Iterate over a list of elements, where these elements have fields that are a
    string representation of a list, such as 'pokemonType' and 'evolutionFamily'.
    And
    that are a string representation of a number such as 'pokemonNumber'.

    '["element_1", "element_2"]' -> ["element_1", "element_2"]
    '1' -> 1

    Arguments:
    - items: list of elements to apply literal_eval on
    """
    for record in items:

```

```

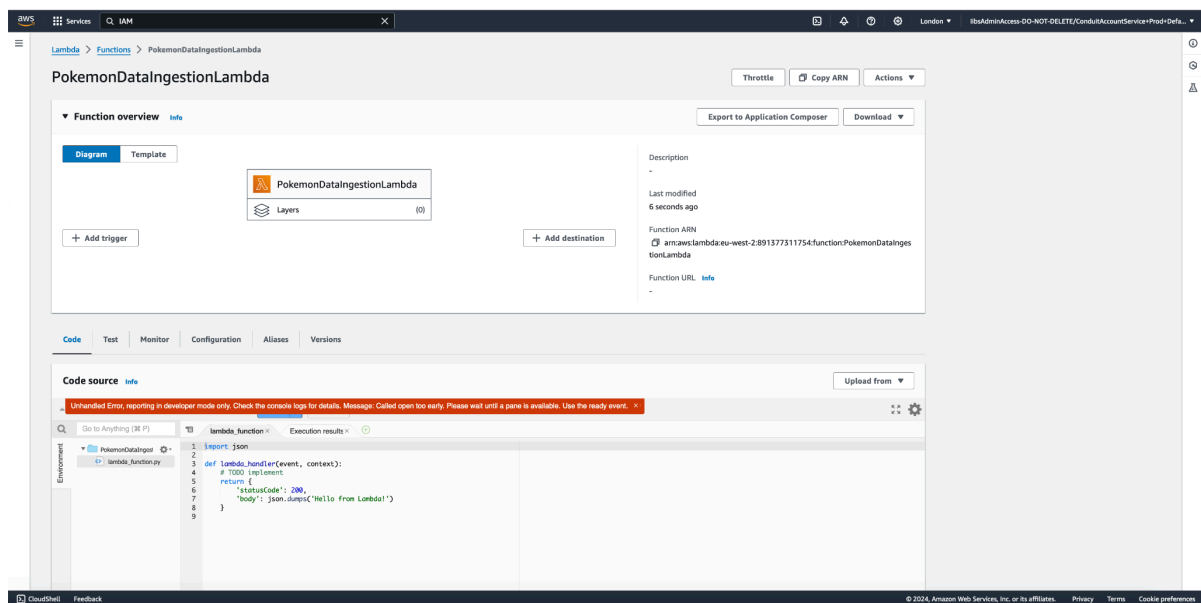
    for key in ["pokemonNumber", "attack", "defense", "evolutionFamily",
               "healthPoints", "speed", "pokemonType"]:
        record[key] = ast.literal_eval(record[key])

def lambda_handler(event, context):
    # Read pokemon dataset
    pokemons_csv_content = read_csv_from_s3_bucket(bucket_name, file_key)
    csv_reader = csv.DictReader(io.StringIO(pokemons_csv_content))
    pokemons = []
    for pokemon in csv_reader:
        pokemons.append(pokemon)

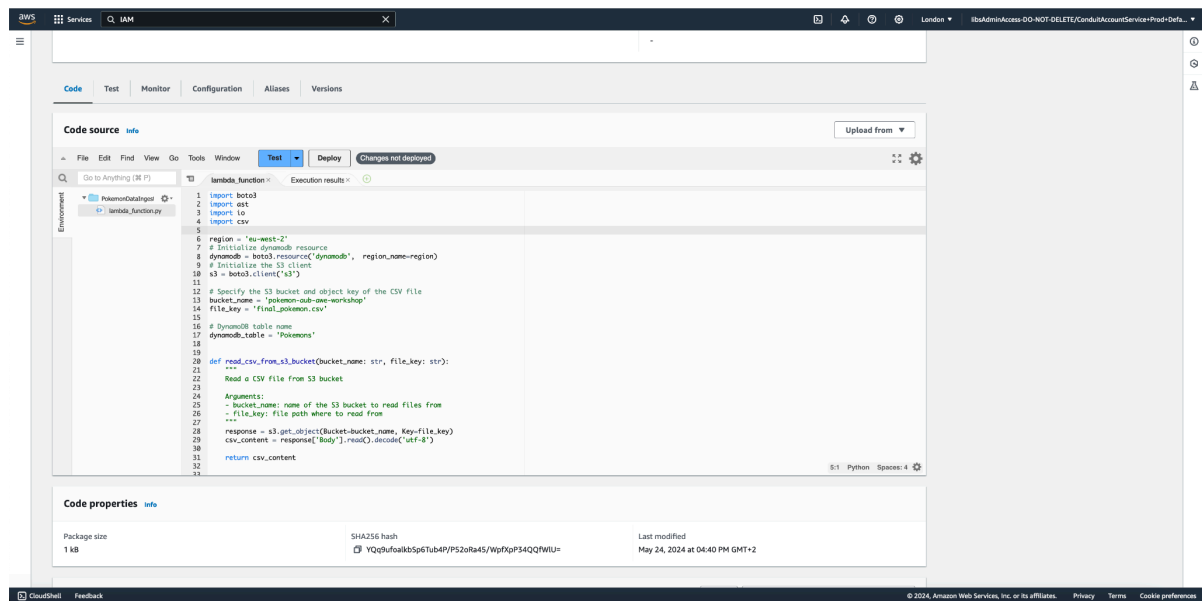
    apply_literal_eval(pokemons)

    # Insert pokemons to DynamoDB, pokemons[1:] because Bulbassaur was manually
    inserted
    insert_dynamoDB_items(dynamodb_table, pokemons[1:])

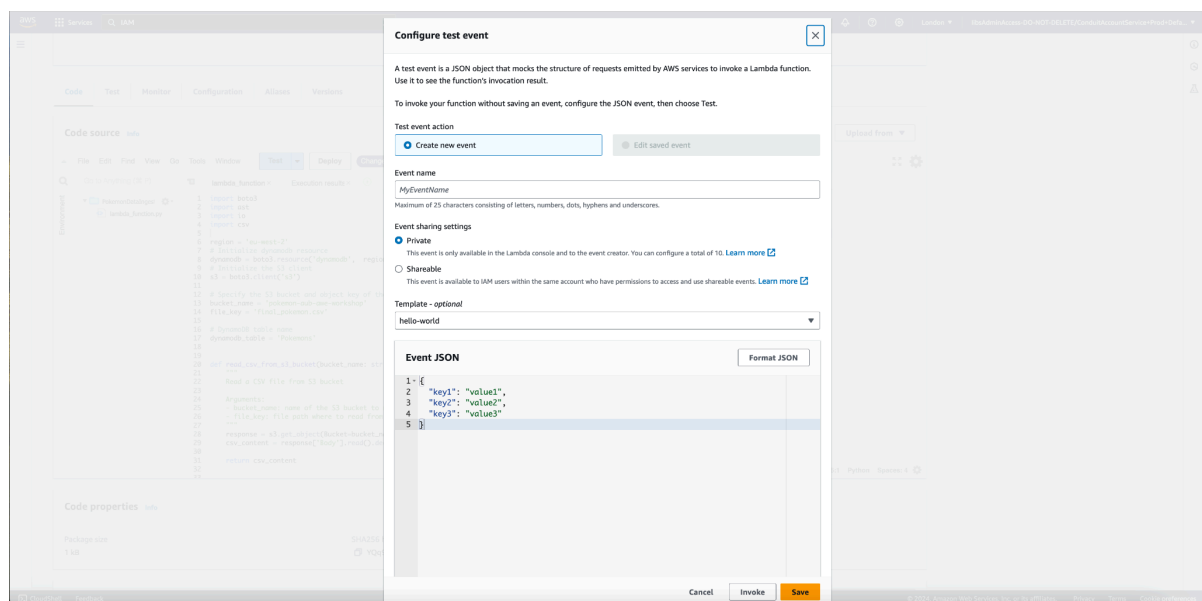
```



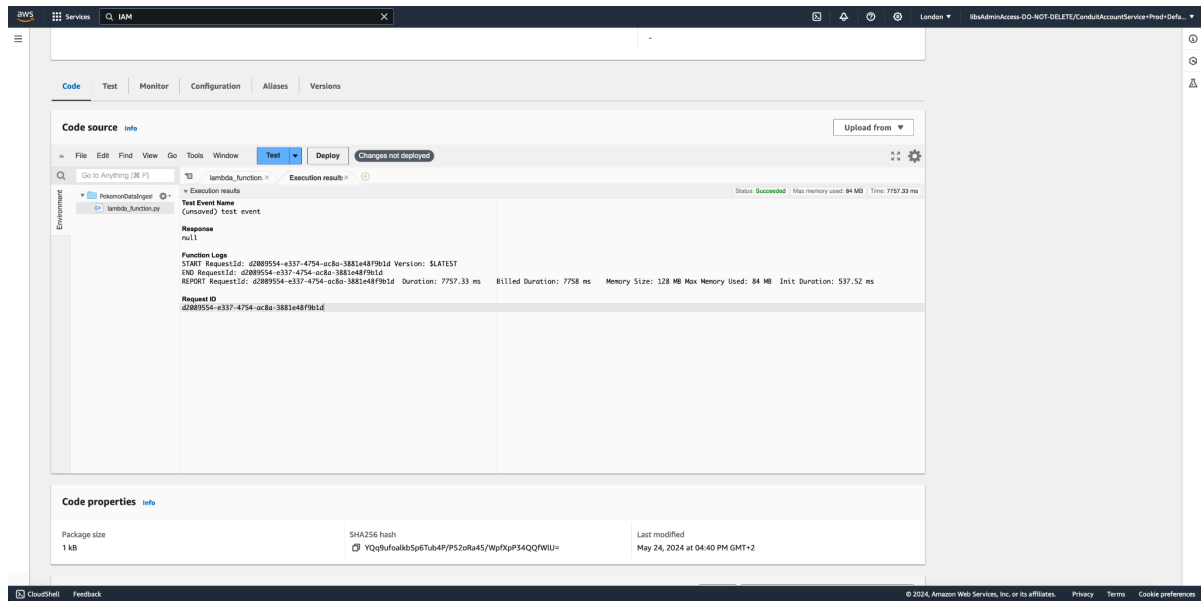
Make sure to **Deploy** your changes, and then **Test** them.



On the Configure test event page, click **Invoke**:

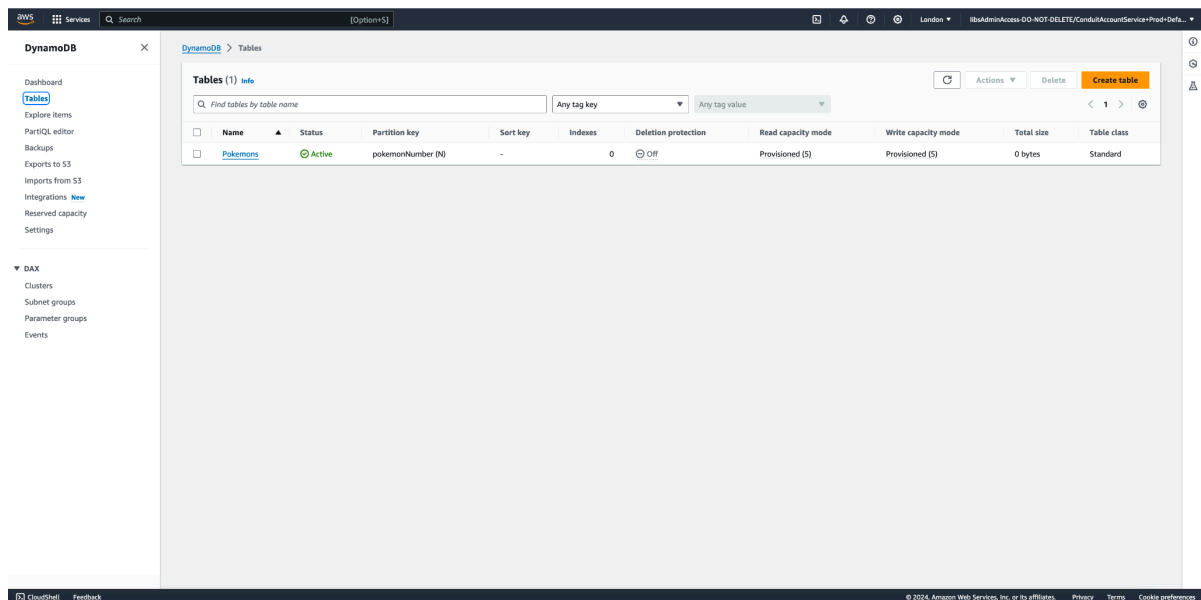


On the **Execution results** tab you should see something similar to the following, but notice that the **requestId** will change every time you execute your lambda.

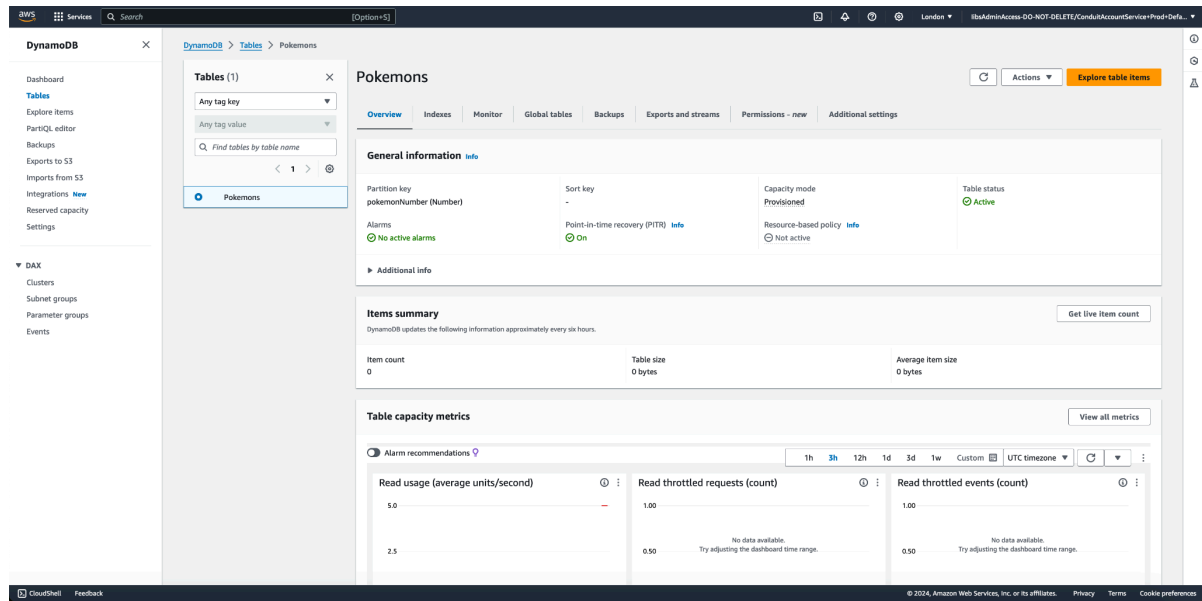


Step 6 - Verify that your Pokemons DynamoDB was correctly populated

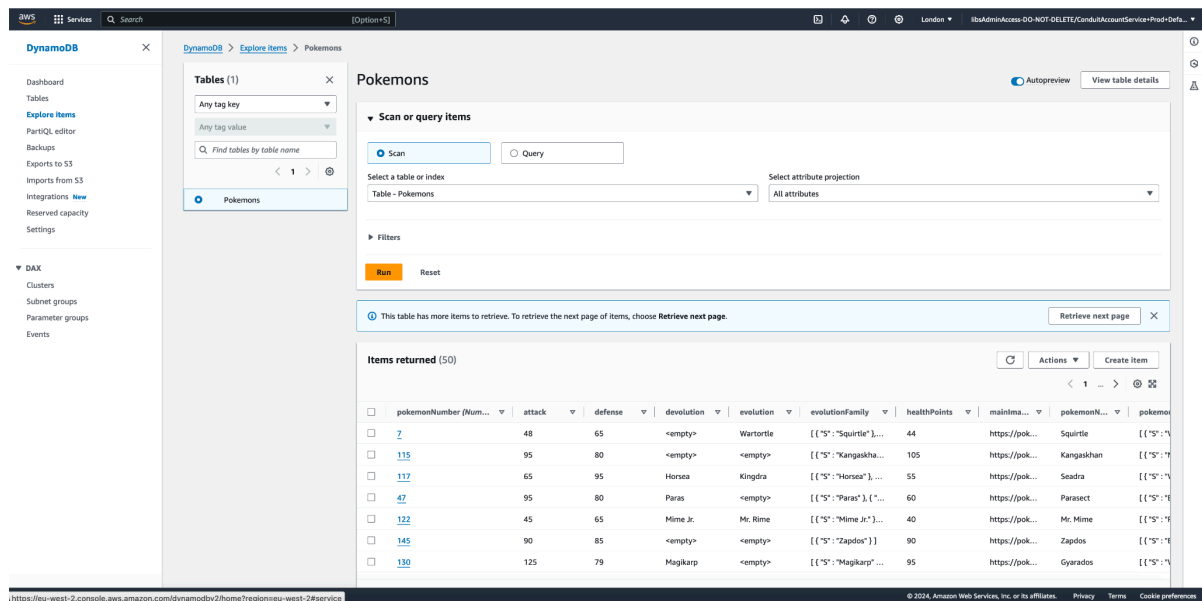
Now, if you go back to your DynamoDB table, you should see that there are 151 pokemons added. Search for DynamoDB again, and then open your Pokemons table:



And select **Explore Table Items**:

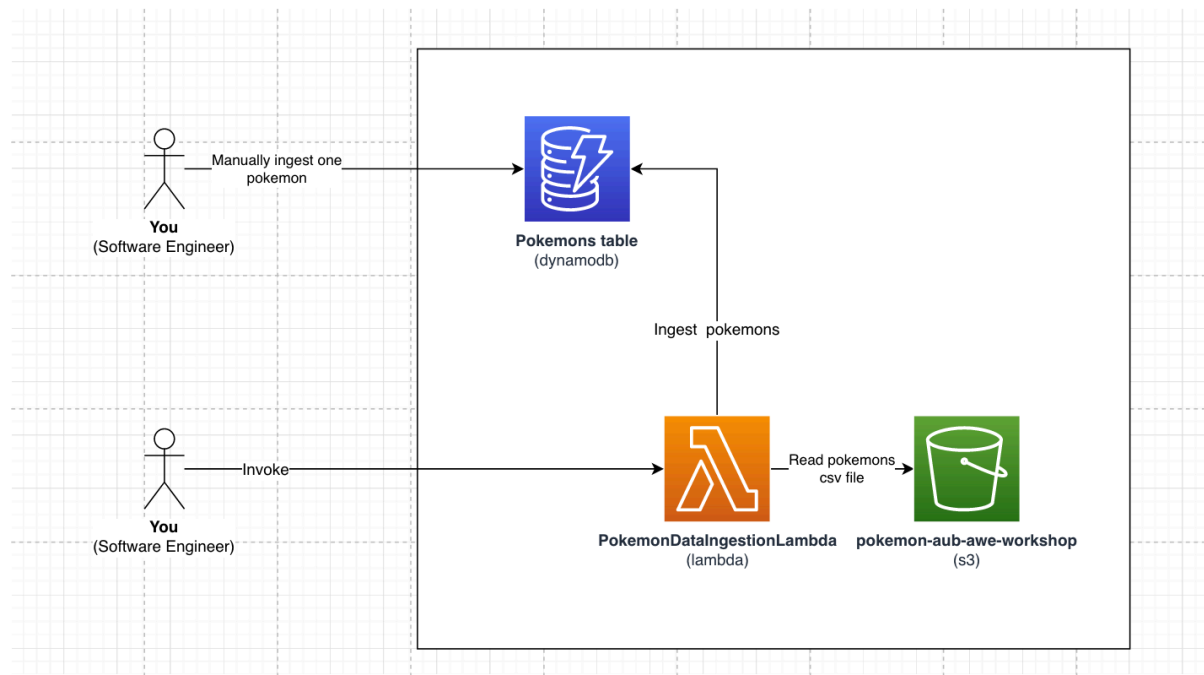


By default, DynamoDB loads the first 50 items from your table, but you can retrieve more using the **Retrieve Next Page** button:



Congratulations 🎉, you have successfully ingested data for all 151 pokemons from the 1st generation of Pokemons to a DynamoDB table. Let's do a quick recap of what we have achieved already:

- Created a DynamoDB table named **Pokemons**
- Ingested one pokemon manually using the AWS console to the **Pokemons** table
- Created a Lambda that ingests pokemons from data from a csv file stored in S3 to the DynamoDB table
- Invoked the Lambda



This DynamoDB table will be the data store used for the application we are building. In the next parts of this workshop, we will be creating an API that will make read requests to this data store; and also create a frontend that will call this API to display pokemons on a website.